

Ronny Alfaro

ENGINEERING LEADER | AI SYSTEMS | PLATFORM ARCHITECTURE | SUPPORT ENGINEERING

+506 6108 5555 | ralfaro@pm.me | [linkedin.com/in/ronniealfaro](https://www.linkedin.com/in/ronniealfaro) | github.com/ronnie-alfaro | [credly.com/users/ralfaro](https://www.credly.com/users/ralfaro)

PROFESSIONAL SUMMARY

Hands-on engineering leader with 20+ years of experience defining technical strategy and solving complex engineering and operational challenges through scalable software platforms, AI-powered systems, cloud-native infrastructure, and developer productivity solutions. Proven record leading global distributed engineering teams while remaining hands-on across software architecture, platform engineering, Python, TypeScript, AI Engineering, Generative AI, Retrieval-Augmented Generation (RAG), Large Language Models (LLMs), automation, DevOps, and cloud technologies. Delivered measurable business impact by increasing a core SLA from ~75% to nearly 99%, building internal engineering platforms supporting over 10,000 customer support requests, driving operational excellence through data-driven optimization, and achieving the highest Individual Contributor engineering level at Experian CheetahMail.

CORE COMPETENCIES

AI & Intelligent Systems

RAG • MCP • AI Agents • Semantic Search • Local LLMs •
LangGraph/LangChain • Vector Databases • PyTorch

Cloud & Reliability

AWS • Azure • GCP • Kubernetes • Terraform • Docker •
CI/CD • Observability • SRE

Platform & Architecture

Platform Engineering • System Design • Distributed
Systems • Developer Tooling • Workflow Automation •
GitOps

Leadership

Engineering Leadership • Technical Strategy • Stakeholder
Management • Mentoring • Incident & Escalation
Management
Agile, Scrum, OKR • RCA

PROFESSIONAL EXPERIENCE

Support Engineering Manager | GitLab

02/2020 – 07/2026

- Led a globally distributed team of 14 engineers across the United States, Canada, Mexico, Brazil, and Costa Rica, supporting ~5,000 customer cases annually and developing engineers into management and Staff-level roles.
- Redesigned workflows and operating processes that raised First Response Time SLA performance from ~75% to nearly 99% through automation, standardization, and Lean Six Sigma methods.
- Architect and built Mechanizer, a Python-based automation platform then turned Zendesk-integrated automation platform used by Support and Product; processed 10,000+ internal and customer requests while reducing common execution time from 15–20 minutes to ~30 seconds. Saving ~500 Engineers hours a month.
- Designed STAR (*Support Ticket Attention Request*), a Zendesk-integrated priority-escalation app adopted by 100+ users across Support, Sales, Solutions Architecture, Professional Services, SRE, and executive leadership.
- Led research, architecture, POCs, presentations, implementation, and model integration for SEMRA, an AI initiative for Support Engineering spanning knowledge retrieval, kb creation, knowledge management and bug issues creation.
- Developed and mentored three engineers who were promoted into Staff Engineer and Engineering Management roles, contributing in improvements in team retention.
- Capacity Planning: Built a data-driven capacity model that proved the business could operate with 15% fewer engineers through workflow optimization rather than increased staffing. Designed and implemented sustainable

operating practices that enabled my team to consistently achieve nearly 2x target performance without increasing workload or creating burnout.

Senior Manager, Global Operations | TradeStation Technologies

04/2017 – 02/2020

- Led ~20 engineers across Systems Administration, Development, ServiceNow, Monitoring, and SREs. supporting latency-sensitive financial trading platforms and multiple server farms of ~100–150 systems.
- Led enterprise adoption of ServiceNow, replacing a constrained in-house platform with automated operational workflows, incident processes, and automated on-call phone escalation.
- Redesigned Linux and Windows patching and operational processes to improve deployment consistency, reliability, and engineering efficiency across production infrastructure.
- Designed Power BI dashboards for IT Operations and the CTO, creating executive visibility into operational performance while mentoring a team member who later transitioned into the Data organization.

Systems Engineer Expert (Staff|Principal Level) | Experian

11/2011 – 04/2017

Career progression: Data Center Operations Engineer Sr → Application Operations → Application Operations Engineer Sr → Team Lead → Systems Engineer Expert

- Progressed through five technical and leadership roles to Systems Engineer Expert, the organization's highest technical IC level, comparable to Staff/Principal Engineer.
- Led an 8-person Operations Service Desk and expanded first-line technical ownership so effectively that L2 Application Operations was ultimately consolidated into Engineering.
- Owned critical domains of CheetahMail, a massive-scale enterprise email platform, including data ingestion, recipient processing, private customer delivery networks, and shared public delivery infrastructure.
- Led design of a dedicated event-driven email delivery network for a major global video game company; provided senior production support across Perl, Oracle Exadata, Linux, networking, and incident response.
- Built Python automation validating NS delegation, MX, SPF, DKIM, reverse DNS, and routing integrity, reducing selected troubleshooting workflows from ~30 minutes to ~5 minutes;
- Led Agile transformation of operations workflows using Scrum ceremonies and OKR-based goal setting.

Senior Solution Architect / Senior Sysadmin | Arcus (Comcast Partner)

09/2010 – 11/2011

- Designed early AWS infrastructure using EC2, ELB, RDS, Route 53, VPC foundations, and early Auto Scaling for a dedicated client.
- Embedded directly within Comcast as an L3 systems engineer supporting large-scale Linux environments; handled security hardening, certificates, audits, migrations.

Head of Engineering | Event by Wire

09/2008 – 09/2010

- Led the 10-person technology organization within a ~30-person startup delivering live internet broadcasting services for funeral homes, owning engineering, architecture, infrastructure, and customer-facing technical operations across the United States.
- Architected and enhanced a Windows Media Services streaming platform, designed secure air-gapped deployment appliances, built multimedia tribute software, and developed an end-to-end digital fulfillment platform for digital and DVD assets.

SELECTED AI / ENGINEERING PROJECTS

Metaphrasis | github.com/ronnie-alfaro/Metaphrasis

- Built a resilient local-first long-form EPUB translation engine using GGUF models, worker-based execution, SQLite-backed state, checkpointing, pause/resume, crash recovery, and automatic LLM-driven self-healing for malformed document structures.
- Preserves literary coherence across English, French, and Spanish; achieves ~70 tokens/sec on an M4 Max and translates a typical 200-page book in ~2 hours without commercial inference APIs.

librerIA | <https://github.com/ronnie-alfaro/librerIA>

- Built a local-first Retrieval-Augmented Generation platform for literary analysis rather than document chat, using semantic EPUB chunking, locally generated embeddings, ChromaDB, SQLite metadata/state, cosine similarity, and MCP.
- Produces character relationship maps, chapter summaries, narrative timelines, key-event highlights, and RPG-style character profiles through a modern analysis UI.

TECHNICAL SKILLS

Languages

Python • TypeScript • JavaScript • Rust • SQL • Bash • Perl

Platform / Cloud

Docker Compose • GitHub Actions • GitOps • Kubernetes • Terraform • CloudFormation • Ansible • AWS • Azure • GCP • Linux

AI / Retrieval

RAG • MCP • LangChain • LangGraph • Embeddings • Semantic Search • Vector Databases • llama.cpp • GGUF • Local Inference

Data / Reliability

PostgreSQL • SQLite • MongoDB • Redis • Kafka • ChromaDB • Qdrant • Grafana • Prometheus • OpenTelemetry • Splunk • Nginx

RECENT PROFESSIONAL DEVELOPMENT

IBM Developer Skills Network: Advanced RAG with Vector Databases and Retrievers • Agentic AI with LangChain and LangGraph • Vector Databases for RAG • Build RAG Applications • Develop Generative AI Applications • Build AI Agents Using MCP